# Teddi Worledge

*PhD candidate working on legal applications, LLM citations, and human-AI collaboration*

worledge@stanford.edu | Stanford, CA | LinkedIn | GitHub | Personal

## EDUCATION

**Stanford University** 2022 - Present

*Computer Science PhD in Artificial Intelligence • Advisor: Carlos Guestrin • Expected Graduation: 2027*

**Teaching:** CS 281: Ethics of Artificial Intelligence • **Community:** NeurIPS ATTRIB 2026 Organizer

**UC Berkeley** 2018 - 2022

*Computer Science B.A. • Research Advisor: Ben Recht*

## EXPERIENCE

**Crosby Intelligence • Member of Technical Staff** Summer 2026

- Trained open weight LLMs to emulate attorney judgments from production legal data. Engineered Crosby's first model training infrastructure, designed experiments, cleaned data and analyzed results.
- Designed a contract redlining benchmark. Designed randomization and UI for expert annotation.

**Stanford Law School Deborah L. Rhode Center • AI Fellow** 2025 - Present

- Built an LLM assistant for the Los Angeles Superior Court (LASC) to recommend judgments in 30k debt collection cases per year with trustworthy citations to case filings. RCT deployed in Sept 2026.
- Ran a study with 65+ legal scholars; the assistant reduced errors by >40% and review time by 25%.
- Worked with LASC attorneys to translate CA Civil Code §§ 1788.58–60 into technical terms and a UI.

**Stanford Artificial Intelligence Lab • PhD Researcher** 2022 - Present

- Conducted human-subject experiments to quantify trade-offs across 6 LLM systems (implemented 5); observed a 50% reduction in sentences fully supported by citations, +200% perceived utility, and 3x longer user verification times as outputs became more abstractive w.r.t. cited sources.
- Developed a unifying framework for LLM citations over training data attribution and NLP citation to enable the design of more reliable and verifiable LLM applications in high-stakes domains.

**Berkeley Artificial Intelligence Research • Research Assistant** 2018 - 2021

- Trained deep neural networks with ML fairness algorithms (e.g., group DRO) and applied scaling laws to project optimal training data composition across population subgroups at larger dataset sizes.

**Dropbox, Core Desktop Team • Software Engineering Intern** Summer 2020

- Engineered and shipped a Python backend for the user-facing sync history. Collaborated with cross-functional partners to deliver an intuitive user experience.

## PUBLICATIONS

**Theodora Worledge,** Othman Bensouda Koraichi, Daniel Bernal, Aviv Caspi, Tatsunori Hashimoto, Carlos Guestrin, and David Freeman Engstrom. "AI Assistance for Court Review of Default Judgments." *AIES.* 2026.

**Theodora Worledge**, Tatsunori Hashimoto, and Carlos Guestrin. "The Extractive-Abstractive Spectrum: Uncovering Verifiability Trade-offs in LLM Generations." *TACL.* 2025.

**Theodora Worledge***, Judy Hanwen Shen*, Nicole Meister, Caleb Winston, and Carlos Guestrin. "Unifying Corroborative and Contributive Attributions in Large Language Models." *SaTML.* 2024. * Co-first authors

Esther Rolf, **Theodora Worledge**, Ben Recht, and Michael Jordan. "Representation Matters: Assessing the Importance of Subgroup Allocations in Training Data." *ICML.* 2021.

## TALKS & AWARDS

**Talks:** LA Superior Court Civil Bench Judicial Seminar, Orrick Labs Summer Sessions, ACM CS & Law, NeurIPS ATTRIB, JURIX AIDA2J, Stanford Legal Design AI+A2J, Code for Humanities

**Fellowships & Grants:** Crosby Intelligence Fellow (2026), NSF Graduate Research Fellow (2022–2027), HAI-Google Research Token (2025), UC Berkeley Regents' & Chancellor's Scholar (2018–2022)